

InterviewAI — Complete Project Technical Documentation (A to Z)


System Type: Multimodal Artificial Intelligence Mock Interview & Diagnostic Platform
Core Technologies: React 18 • Vite • FastAPI • Qwen3 LLM • Whisper-v3 • DeepFace • Web Audio API • MongoDB
Author: Sumit Kumar Sahu & Team | **Release:** v3.2.0 | **Status:** Production Ready

1. Executive Summary & Problem Statement

InterviewAI is an end-to-end, multi-modal automated mock interview platform engineered to solve the placement readiness crisis faced by software engineering students and early-career candidates. Traditional preparation methods suffer from massive limitations: human mock interviews are expensive (Rs 2,000 - 5,000 per session), standard chatbots provide superficial text answers without verbal interaction, and generic practice tools do not inspect the candidate's actual resume or non-verbal communication.

Traditional Interview Preparation	InterviewAI Multimodal Platform
- Asks generic icebreakers ('Tell me about yourself'). - Cannot cross-examine technical project trade-offs. - Ignores vocal speed (WPM), pauses, and filler words. - No eye contact tracking or emotional sentiment cues. - High commercial fees; inaccessible to regular students. - Static question sets that do not adapt to student level.	- Directly opens with candidate's actual resume project. - Listens to spoken audio and generates dynamic counter-questions. - Vectorized audio analytics for WPM, pauses, and fillers. - DeepFace & OpenCV real-time gaze and posture feedback. 100% self-contained, automated, and free for students. - Full-duplex WebSocket real-time conversational pipeline.

2. High-Level System Architecture & Flow



End-to-End User Journey (A to Z):

- 1. Authentication & Session Initialization:** Candidate signs up/logs in with bcrypt hashed credentials; receives a signed JWT bearer token stored securely in localStorage.
- 2. Resume Ingestion & ATS Parsing:** Candidate uploads a PDF/DOCX resume. The backend extracts clean text, identifies technical skills (languages, frameworks, databases, tools), extracts academic projects, and computes a 0-100 ATS compatibility score.
- 3. Candidate Command Center:** Displays real-time AI service health, personalized greeting, 4 instant practice tracks, today's Rapid Drill question with scratchpad, and a pre-flight hardware calibrator.
- 4. Pre-Flight Hardware Calibration:** Directly inside the dashboard, candidate tests their microphone live. HTML5 Web Audio API calculates real-time RMS amplitude and provides visual feedback before entering the interview.
- 5. Live Interview Room (WebSocket or REST):** Web Audio API records voice in lossless 16kHz 16-bit Mono WAV. OpenCV tracks candidate eye contact and facial sentiment. Whisper transcribes speech, and Qwen3 LLM generates conversational counter-questions.
- 6. Multimodal Diagnostic Report & PDF:** Generates comprehensive scorecards, radar charts, grammar feedback, speech pace (WPM), and downloadable PDF certificate.

3. Detailed Module Breakdown & Algorithms

Module 1: Resume Context Ingestion & ATS Scoring Engine

Files: backend/app/services/resume_parser.py, backend/app/services/ats_scorer.py

- **Format Ingestion:** Supports PDF, DOCX, and plain TXT files. Extracts text streams using pypdf and python-docx.
- **Entity Extraction:** Employs regex heuristic parsers to extract candidate metadata into structured JSON: candidate_name, skills (categorized into languages, frameworks, databases, developer tools), projects (titles, tech stacks, bullet points), and education.
- **ATS Scoring Formula:** Evaluates 3 weighted components:

$$\text{Score} = 0.50 \times (\text{Keyword Match Ratio}) + 0.30 \times (\text{Section Completeness}) + 0.20 \times (\text{Formatting \& Density})$$

Outputs matched skills, missing target keywords for the selected role, and actionable bullet-point tips.

Module 2: Real-Time Dynamic Conversational AI Engine

Files: backend/app/services/interview_engine.py, backend/app/services/prompt_templates.py

- **Context-Grounded Opening:** Instead of asking 'Tell me about yourself', the system immediately probes: *'I see on your resume that you built [Top Project Name] using [Tech Stack]. Can you explain the architectural decision behind that?'*
- **Dynamic Counter-Questioning:** Analyzes the candidate's actual answer text. If the candidate gives a good high-level answer, the LLM immediately asks about implementation edge cases or testing. If the candidate is stuck, it warmly offers a foundational hint.
- **Strict Student Governance:** Built-in guardrails forbid asking senior enterprise scalability dilemmas (e.g. 500,000 QPS Kafka pipelines) and keep questions grounded in college-level projects, data structures, and standard REST APIs.
- **Resilient Fallback Hierarchy:** Primary model: Qwen/Qwen3-8B; Fallback 1: Qwen/Qwen3-4B; Fallback 2: Local Ollama Qwen2.5:7b; Fallback 3: Rule-based contextual mock engine.

Module 3: In-Browser Lossless Audio Encoding & Speech-to-Text

Files: backend/app/services/speech_service.py, frontend/src/pages/InterviewRoom.jsx

- **Browser Audio Incompatibility Solved:** Standard browser MediaRecorder outputs WebM/Opus audio, which libsndfile on HuggingFace rejects with a 400 Bad Request error. We implemented a client-side Web Audio API PCM decoder that encodes recorded audio directly into standard **16kHz 16-bit Mono RIFF WAVE**. This guarantees 100% transcription reliability across all browsers including Brave and Safari.
- **Whisper-v3 Inference:** Transcribes audio into clean text with punctuation, capitalization, and latency under 1.2 seconds.

Module 4: Ultra-Fast Vectorized Voice Analytics (170s → 0.02s)

Files: backend/app/services/audio_service.py

- **Optimization Breakthrough:** Traditional audio packages like Librosa execute nested Fourier transform loops synchronously, freezing the server CPU at 88% for 170+ seconds on an 8-second clip. We engineered a custom vectorized algorithm using **Soundfile + NumPy root-mean-square (RMS) energy thresholding**.
- **Execution Time:** Drops from **172.4 seconds to 0.02 seconds** (over 8,500x speedup) with zero event-loop blocking.
- **Extracted Vocal Metrics:** Total audio duration, speaking duration, pause duration percentage, speaking rate in Words Per Minute (WPM), and filler word frequencies ('um', 'uh', 'like', 'you know').

Module 5: Computer Vision & Non-Verbal Sentiment Tracking

Files: backend/app/services/cv_service.py, frontend/src/components/CameraPanel.jsx

- **Eye Contact Vector:** Uses client-side HTML5 Canvas face landmark bounding box to calculate distance from optical center. Translates into a 0-100% eye contact stability score.
- **Facial Emotion Analysis:** Uses DeepFace to classify emotional distribution across frames (Neutral, Happy, Confident, Stressed) to assist the candidate with composure feedback.

Module 6: Candidate Command Center & Dashboard Upgrades

Files: frontend/src/pages/Dashboard.jsx, AIStatusPanel.jsx, DailyChallengeCard.jsx, HardwareCalibrator.jsx

- **Compact AI Status Ribbon:** Slim 44px glassmorphic strip showing real-time live pulse for LLM, Whisper, Vision, and Semantic engines with collapsible drawer.
- **Instant Practice Tracks:** 4 one-click launch cards: Technical & System Design, Behavioral (STAR Method), Resume Deep-Dive, and HR Screening.
- **Daily Rapid Drill:** Features rotating daily interview challenges, candidate scratchpad notes, streak tracking (Day 1 Streak ■), and expandable model answer frameworks.
- **Pre-Flight Hardware Calibrator:** Live microphone volume level meter using Web Audio API and real-time backend ping latency tracker.
- **Interview Readiness Roadmap:** 3-stage milestone tracker (Resume ATS Indexing → First Diagnostic Session → Speech & Facial Review).

Module 7: Global Alert & Notification System

Files: frontend/src/context/AlertContext.jsx

- Replaces all disruptive browser alert() and window.confirm() popups with a modern, glassmorphic toast & confirmation dialog suite.
- Stackable toasts for **Success** (emerald), **Error** (crimson), **Warning** (amber), and **Info** (purple) with auto-dismiss progress timers.
- Promise-based interactive confirmation modals for session deletion, resume deletion, and logout.

4. Complete API & WebSocket Catalog

Method & Endpoint	Auth	Request / Query	Response & Purpose
POST /api/auth/register	None	email, password, first_name, last_name	Creates user in MongoDB, returns JWT token.
POST /api/auth/login		email, password (form-data)	Validates password, issues JWT bearer token.
GET /api/auth/me	Bearer	None (extracts from JWT)	Returns current authenticated user profile.
POST /api/resume/upload	Bearer	Multipart: file (PDF/DOCX)	Parses text, extracts skills & projects, scores ATS.
GET /api/resume	Bearer	None	Returns candidate's active resume & ATS score.
DELETE /api/resume	Bearer	None	Removes resume profile and resets customized skills.
POST /api/interviews	Bearer	role, experience_level, interview_type	Initializes session, sets resume opening question.
GET /api/interviews	Bearer	statusFilter, searchQuery	Lists all previous interview sessions and scores.
GET /api/interviews/{id}	Bearer	Path: id	Fetches full transcript, scores, and questions.
POST /api/interviews/{id}/answer	Bearer	Multipart: audio (16kHz WAV) or text	Transcribes, evaluates, generates counter-question.
POST /api/interviews/{id}/complete	Bearer	Path: id	Computes final scores and creates report debrief.
GET /api/reports/{id}/pdf	Bearer	Path: id	Downloads compiled PDF evaluation certificate.
GET /api/ai/health	None	None	Checks status of LLM, Whisper, CV, and Embeddings.
WS /ws/interview/{id}	Ticket	Full-duplex WebSocket stream	Streams real-time question/answer/TTS/evaluation.

5. Complete Database Schema (MongoDB Collections)

```
1. users Collection
{ _id: ObjectId, email: String (unique), first_name: String, last_name: String,
  hashed_password: String, created_at: ISODate, updated_at: ISODate }

2. resumes Collection
{ _id: ObjectId, user_id: ObjectId (indexed), filename: String, parsed_text: String,
  skills: [String], experience: [String], education: [String],
  resume_context: {
    candidate_name: String,
    skills: { languages: [String], frameworks: [String], databases: [String], tools: [String] },
    projects: [ { title: String, tech_stack: [String], description: String } ]
  },
  ats_score: Number (0-100), ats_summary: String,
  ats_matched_keywords: [String], ats_missing_keywords: [String], uploaded_at: ISODate }

3. interviews Collection
{ _id: ObjectId, user_id: ObjectId (indexed), role: String, experience_level: String,
  interview_type: 'Technical'|'Behavioral'|'Resume-Based'|'HR'|'Mixed',
  status: 'active'|'completed', overall_score: Number (0-100), duration_seconds: Number,
  scores_breakdown: {
    technical: Number, communication: Number, confidence: Number,
    eye_contact: Number, speech_clarity: Number, answer_quality: Number
  },
  feedback: {
```

```
    strengths: [String], improvements: [String], summary: String,
    pace_feedback: String, filler_word_feedback: String
  },
  questions: [
    {
      id: Number, question_text: String, answer_text: String, audio_path: String,
      eye_contact_score: Number, emotion_summary: Object, voice_metrics: Object,
      evaluation: { score: Number, strengths: [String], improvements: [String], feedback: String }
    }
  ],
  created_at: ISODate, completed_at: ISODate }
```

6. Project Viva Voce & Technical Defense Q&A;

Q1: What is the core innovation of InterviewAI compared to commercial tools?

Technical Answer: Commercial tools like Interviewing.io or Pramp either require pairing with human engineers (expensive and scheduling-dependent) or rely on generic chatbot interfaces that ignore vocal delivery and body language. InterviewAI unites three distinct modalities: (1) Resume-grounded conversational LLM counter-questioning, (2) Vectorized vocal pacing and filler word analytics, and (3) Computer vision gaze and sentiment tracking, running in real-time on standard consumer hardware.

Q2: Why did you implement client-side Web Audio API WAV encoding instead of standard browser MediaRecorder?

Technical Answer: Standard browser MediaRecorder captures audio in WebM/Opus format. While WebM is compact, backend inference services (like Hugging Face Inference API and Soundfile) depend on libsndfile, which does not natively decode WebM container packets and throws HTTP 400 Bad Request. By decoding Float32 PCM directly in the browser and formatting clean 16kHz 16-bit Mono WAV headers, we eliminated all transcoding overhead and achieved 100% audio compatibility across Chrome, Safari, and Brave.

Q3: How did you fix the severe backend freezing issue during audio analysis?

Technical Answer: Originally, the backend utilized Librosa's piptrack and harmonic analysis algorithms, which executed nested mathematical loops synchronously on the main FastAPI event loop, spiking CPU to 88% and stalling requests for 170+ seconds. We replaced it with vectorized Soundfile and NumPy RMS energy thresholding. This reduced processing time from 172.4 seconds down to 0.02 seconds—an 8,500x acceleration—allowing real-time audio debriefs without blocking async requests.

Q4: How does the dynamic counter-questioning state machine work?

Technical Answer: The conversational engine maintains state tracking the current topic, consecutive turns, and previous responses. When an answer is received, Qwen3 LLM analyzes if key technical concepts were articulated. If the candidate answers well, the model reinforces the point and asks an edge-case implementation or testing question. If the candidate expresses confusion ('I don't know' or hesitates), the engine warmly pivots to a foundational concept to build confidence.

Q5: How does the ATS Resume Scoring engine evaluate candidates?

Technical Answer: The ATS engine extracts structured entities and computes a composite score based on: (1) Technical keyword match against standard industry skills for the selected role (50% weight), (2) Structural completeness verifying dedicated Education, Experience, Projects, and Skills sections (30% weight), and (3) Formatting hygiene and keyword density (20% weight). It also returns a missing skills checklist.

Q6: How does the Pre-Flight Hardware Calibrator assist candidates before interviews?

Technical Answer: Hardware anxiety is a primary reason candidates fail online interviews. The Pre-Flight Calibrator uses the browser's Web Audio API AudioContext and AnalyserNode to calculate real-time microphone RMS volume levels, displaying an interactive decibel meter. It also confirms webcam readiness and pings the backend server to display real-time WebSocket speech pipeline latency (~45 ms).

Q7: How is candidate privacy and ethical AI ensured in video/audio processing?

Technical Answer: All audio and webcam processing can operate on local or zero-retention pipelines. Video frames for eye contact are evaluated locally in the browser via HTML5 Canvas math or transiently passed to DeepFace without storing raw video streams. A permanent disclaimer informs candidates that facial sentiment metrics are practice aids for posture, not psychological evaluations.

Q8: What fallback strategies are in place if the primary LLM API goes down?

Technical Answer: The backend uses a 4-tier resilient fallback hierarchy: Primary is Qwen3-8B via Hugging Face Router; if rate-limited or unavailable, it automatically falls back to Qwen3-4B; if external internet is constrained, it routes to a local Ollama Qwen2.5 instance; and if offline, it uses an internal rule-based technical question generator, ensuring zero session disruption.

Q9: How did you design the notification system to improve user experience?

Technical Answer: We replaced disruptive native browser alert() and confirm() popups with a centralized AlertContext. It provides non-blocking, stackable glassmorphic toasts with auto-dismiss countdown bars for Success, Error, Warning, and Info, plus interactive confirmation dialogs for destructive actions like session deletion, resume deletion, and logout.

Q10: What are the future enhancement opportunities for InterviewAI?

Technical Answer: Future milestones include: (1) Integrating full WebRTC data channels for sub-100ms conversational audio streaming, (2) Adding live collaborative code execution sandboxes for coding rounds, (3) Implementing multi-interviewer panel simulations (Technical Lead + HR Manager), and (4) Exporting standardized candidate readiness badges for LinkedIn and campus placement portals.